

COMP 251: Algorithms & Data Structures

Owen Lewis

Winter 2018

Contents

1	Overview of Graph Theory	2
1.1	Definitions	2
1.2	Some Theorems for Undirected Graphs	2
1.3	““Data Structures”” for Representing Graphs	4
1.3.1	Adjacency Matrices	4
1.3.2	Adjacency Lists	4
1.3.3	Adjacency Matrices vs. Adjacency Lists	5
2	Divide & Conquer	6
2.0.1	The Master Theorem	6
2.1	MergeSort	8
2.2	Binary Search	9
2.3	Fast Multiplication	10
2.4	Fast Matrix Multiplication	10
2.5	The Selection Problem	10
3	Graph Algorithms	13
4	Greedy Algorithms	14
5	Dynamic Programming	15
6	Network Flows	16
7	Data Structures	17

1 Overview of Graph Theory

1.1 Definitions

A graph $G = (V, E)$ is a set V of vertices (a.k.a. nodes) and a set E of edges (denoting vertex pairs). We set $n = |V|$, and $m = |E|$. A graph is said to be *undirected* when for any edge $(u, v) \in E$ there exists an edge $(v, u) \in E$ for some nodes u , and v . A graph is said to be *directed* if it is not undirected. In other words, the edge set of a directed graph consists of ordered pairs where the edge set of an undirected graph consists of unordered pairs.

A *walk* is a set of vertices $\{v_0, v_1, \dots, v_\ell\}$ such that $(v_i, v_{i+1}) \in E, \forall 0 \leq i \leq \ell$. A walk where $v_0 = v_\ell$ is said to be a *circuit* or a *closed walk*. A circuit where every edge in the graph is used exactly once is known as an *Eulerian circuit*. A *cycle* is a walk $\{v_0, v_1, \dots, v_\ell\}$ such that every vertex is distinct except $v_0 = v_\ell$. A cycle where every vertex of the graph is used exactly once is known as a *Hamiltonian cycle*. A walk where every vertex is distinct is said to be a *path*.

A graph is said to be *connected* if for each $u, v \in V$ there exists a walk from u to v . A graph is said to be *disconnected* if it is not connected. Each connected subgraph of a graph is called a *component*. A connected graph therefore has exactly one component.

A connected component with no cycles is called a *tree*. A graph whose components are all trees is said to be a *forest*. A tree is said to be *spanning* if it contains every vertex in the graph. A vertex in a tree with at most one neighbour is called a *leaf*.

A *matching* is a set of vertex-disjoint edges i.e. each edge is incident to at most one other edge in a matching. A matching is said to be *perfect* if every vertex is incident to exactly one edge in the matching.

A *clique* is a set of pairwise adjacent vertices. In *independent set* (a.k.a. a *stable set*) is set of pairwise non-adjacent vertices.

A *bipartite graph* is a graph such that the vertex set V can be partitioned as $V = X \cup Y$ where each edge has one node in X and the other node in Y . Note that X and Y are necessarily independent sets.

1.2 Some Theorems for Undirected Graphs

Theorem: (Handshaking Lemma) Let $G = (V, E)$ be an undirected graph, let $\Gamma(v) := \{u : (u, v) \in E\}$ be the set of neighbours of a node v , and let the *degree* $\deg(v)$ of a vertex v equal the cardinality of $\Gamma(v)$. Then there are an even number of vertices with odd degree.

Proof: First note that since we're double-counting the number of pairs where

(v, e) is an edge incident to v

$$2 \cdot |E| = \sum_{v \in V} \deg(v)$$

Since the degree of a vertex is either even or odd, we can partition V into a set of odd-degree vertices \mathcal{O} , and a set of even-degree vertices \mathcal{E} . This gives us

$$\sum_{v \in V} \deg(v) = \sum_{v \in \mathcal{O}} \deg(v) + \sum_{v \in \mathcal{E}} \deg(v)$$

which implies

$$\sum_{v \in \mathcal{O}} \deg(v) = 2 \cdot |E| - \sum_{v \in \mathcal{E}} \deg(v)$$

since both the $2 \cdot |E|$ term is even (obvious) and the $\sum_{v \in \mathcal{E}} \deg(v)$ term is even (sum of even numbers) then the $\sum_{v \in \mathcal{O}} \deg(v)$ term must also be even. \square

Theorem: (Euler's Theorem) If G is an undirected graph then G contains an Eulerian circuit if and only if every vertex has even degree.

Proof: Easy proof by induction

Lemma: A tree T with $n \geq 2$ vertices has at least one leaf vertex.

Proof: Trees are connected so there exists no vertices with degree 0 when $n \geq 2$. Suppose each vertex has degree of at least 2. Then consider the longest path $P \subseteq T$, $P = \{v_1, v_2, \dots, v_{\ell-1}, v_{\ell}\}$. Since $\deg(v_{\ell}) \geq 2$, \exists a neighbour (of v_{ℓ}) $x \in P$ with $x \neq v_{\ell-1}$. If $x = v_{\ell+1}$ then P is not the longest path, a contradiction. Therefore, for P to be the longest path, x must be somewhere else in P , but this creates a cycle, another contradiction. Thus there must exist at least one node v such that $0 < \deg(v) < 2$ – a leaf. \square

Theorem: A tree with n vertices has exactly $n - 1$ edges.

Proof: Simple proof by induction.

Base case: A tree with one vertex trivially has 0 edges.

Induction Hypothesis: Assume any tree with $n - 1$ vertices has $n - 2$ edges.

Inductive Step: Take a tree with $n \geq 2$ vertices. By the previous lemma this tree contains a leaf vertex v . This implies that $T \setminus \{v\}$ is a tree with $n - 1$ vertices and by the induction hypothesis $T \setminus \{v\}$ is a tree with $n - 2$ edges, which implies that T is a tree with $n - 1$ edges. \square

Theorem: (Hall's Theorem) Let $G = (X \cup Y, E)$ with $|X| = |Y|$ be a bipartite graph. G contains a perfect matching if and only if $\forall B \subseteq X$, $|\Gamma(B)| \geq |B|$ (Hall's condition).

Proof: Firstly, the (\Rightarrow) direction is fairly obvious. If $B \subseteq X$ with $|\Gamma(B)| < |B|$ then the graph can't have a perfect matching. The (\Leftarrow) direction is a bit trickier. Suppose Hall's condition is satisfied. Then, take the maximum cardinality

matching M is the graph. If M is perfect then we are done. Otherwise there must exist an unmatched vertex b_0 .

- Since Hall's condition holds, we have $|\Gamma(\{b_0\})| \geq |\{b_0\}| = 1$ so b_0 must have at least one neighbour s_0 .
- Suppose s_0 is matched in M to b_1 .
- Since Hall's condition holds, we have $|\Gamma(\{b_0, b_1\})| \geq |\{b_0, b_1\}| = 2$ so $\{b_0, b_1\}$ must have at least one neighbour $s_1 \neq s_0$.
- Suppose s_1 is matched in M to b_2 .
- Since Hall's condition holds, we have $|\Gamma(\{b_0, b_1, b_2\})| \geq |\{b_0, b_1, b_2\}| = 3$ so $\{b_0, b_1, b_2\}$ must have at least one neighbour $s_2 \notin \{s_0, s_1\}$.
- ...

we repeat this argument as long as we can. Since the graph contains a finite number of vertices this process must terminate, but it can only terminate when we reach an unmatched node s_k . Using the edges we've formed in M we can create a path P from b_0 to s_k that alternates between using non-matching edges and using matching edges. Swapping the matching edges with the non-matching edges gives us one more matching edge (as we have an odd number of edges.) This is still a valid matching as the internal nodes of P are still incident to exactly one matching edge. Also, the end nodes, b_0 and s_k were previously unmatched but are now incident to exactly one edge in the new matching. Thus M isn't the maximum capacity matching – a contradiction. \square

1.3 “““Data Structures””” for Representing Graphs

1.3.1 Adjacency Matrices

For a graph, an *adjacency matrix* M is a matrix such that

1. There is a row for each vertex
2. There is a column for each vertex
3. The ij – th entry is defined as $M_{ij} = \begin{cases} 1, (i, j) \in E \\ 0, (i, j) \notin E \end{cases}$

Note that in an undirected graph the matrix is symmetric around the diagonal because $(i, j) \equiv (j, i)$. Of course this is not necessarily true of directed graphs.

1.3.2 Adjacency Lists

An *adjacency list* of an undirected graph is such that for each vertex v of V we store a list of its neighbours. For a directed graph we have two lists: one in which we store the in-neighbours of v and one in which we store the out-neighbours of v .

1.3.3 Adjacency Matrices vs. Adjacency Lists

The main difference between the two is the amount of storage required to implement them.

- An adjacency matrix requires we store $\Theta(n^2)$ numbers
- An adjacency list requires we store $\Theta(m)$ numbers

In any graph $m = O(n^2)$. This means that for a sparse graph adjacency lists are highly favourable in terms of space complexity.

Verifying whether an edge exists, however, is much faster in an adjacency matrix – when using the array representation of a matrix it takes $O(1)$ time, where verifying the existence of an edge takes $O(\log n)$ time for an ordered adjacency list (using binary search), and $O(n)$ time if the adjacency list is not ordered (using sequential search).

2 Divide & Conquer

A *divide and conquer* algorithm ideally breaks up a problem of size n into smaller sub-problems such that:

- There are exactly a sub-problems
- Each sub-problem has a size of at most $\frac{1}{b} \cdot n$
- Once solved, the solutions to the sub-problems must be combined in $O(n^d)$ time to produce a solution to the original problem

Therefore the time-complexity of a divide and conquer algorithm satisfies a recurrence relation given by

$$T(n) = a \cdot T\left(\frac{n}{b}\right) + O(n^d)$$

2.0.1 The Master Theorem

Lemma 1: $\sum_{k=0}^{\ell} \tau^k = \frac{1 - \tau^{\ell+1}}{1 - \tau}$, for any $\tau \neq 1$.

Proof:

$$\begin{aligned} (1 - \tau) \sum_{k=0}^{\ell} \tau^k &= \sum_{k=0}^{\ell} \tau^k - \sum_{k=1}^{\ell+1} \tau^k \\ &= \tau^0 - \tau^{\ell+1} \\ &= 1 - \tau^{\ell+1} \\ \sum_{k=0}^{\ell} \tau^k &= \frac{1 - \tau^{\ell+1}}{1 - \tau} \end{aligned}$$

□

Lemma 2: $x^{\log_b y} = y^{\log_b x}$ for any base $b \in \mathbb{R}$.

Proof: From the power rule of logarithms we have

$$\log_b x \cdot \log_b y = \log_b (y^{\log_b x})$$

similarly, we have

$$\log_b x \cdot \log_b y = \log_b (x^{\log_b y})$$

therefore

$$\log_b (x^{\log_b y}) = \log_b (y^{\log_b x})$$

thus

$$x^{\log_b y} = y^{\log_b x}$$

□

Theorem: (The Master Theorem) If a recurrence relation is of the form $T(n) = a \cdot T\left(\frac{n}{b}\right) + O(n^d)$, for constants $a > 0$, $b > 1$, and $d \geq 0$, then

$$T(n) = \begin{cases} O(n^d) & , \text{ if } a < b^d \text{ [Case I]} \\ O(n^d \cdot \log n) & , \text{ if } a = b^d \text{ [Case II]} \\ O(n^{\log_b a}) & , \text{ if } a > b^d \text{ [Case III]} \end{cases}$$

Proof: By adding dummy numbers, we may assume that n is a power of b , i.e. $n = b^\ell$, for some $\ell \in \mathbb{N}_0$. Then

$$T(n) = n^d + a\left(\frac{n}{b}\right)^d + a^2\left(\frac{n}{b^2}\right)^d + \cdots + a^\ell\left(\frac{n}{b^\ell}\right)^d$$

simplifying, we get

$$T(n) = n^d \left(1 + \frac{a}{b^d} + \left(\frac{a}{b^d}\right)^2 + \cdots + \left(\frac{a}{b^d}\right)^\ell \right)$$

we now have three cases:

Case I: $\frac{a}{b^d} < 1$

Set $\tau := \frac{a}{b^d}$ then we have

$$T(n) = n^d \sum_{k=0}^{\ell} \tau^k$$

applying lemma 1 we get

$$T(n) = n^d \left(\frac{1 - \tau^{\ell+1}}{1 - \tau} \right) \leq n^d \left(\frac{1}{1 - \tau} \right)$$

but $\frac{1}{1 - \tau}$ is a constant so $T(n) \leq n^d$ and thus

$$T(n) = O(n^d)$$

Case II: $\frac{a}{b^d} = 1$

Then we have that

$$T(n) = n^d (1 + 1 + 1^2 + \cdots + 1^\ell) = n^d (\ell + 1)$$

but $n = b^\ell \Rightarrow \ell = \log_b n$, thus

$$T(n) = O(n^d \cdot \log n)$$

Case III: $\frac{a}{b^d} > 1$

Set $\tau := \frac{a}{b^d}$ then we have

$$T(n) = n^d \sum_{k=0}^{\ell} \tau^k$$

applying lemma 1 we get

$$T(n) = n^d \left(\frac{\tau^{\ell+1} - 1}{\tau - 1} \right) \leq n^d \left(\frac{\tau^{\ell+1}}{\tau - 1} \right)$$

but since $\tau - 1$ is a constant we get

$$\begin{aligned} T(n) &= n^d O(\tau^{\ell+1}) \\ &= O(n^d \tau^{\ell+1}) \\ &= O(n^d \tau^{\ell}) \\ &= O\left(\left(\frac{a}{b^d}\right)^{\ell} n^d\right) \\ &= O\left(\left(\frac{n}{b^{\ell}}\right)^d a^{\ell}\right) \end{aligned}$$

but $n = b^{\ell}$ so

$$T(n) = O(a^{\ell})$$

and $\ell = \log_b a$ so

$$T(n) = O(a^{\log_b n})$$

and applying lemma 2 gives

$$T(n) = O(n^{\log_b a})$$

This completes the proof. □

2.1 MergeSort

MergeSort is an algorithm to sort n numbers into non-decreasing order.

```

MergeSort( $x_1, x_2, \dots, x_n$ )
  if  $n = 1$ 
    return  $x_1$ 
  else
    return Merge(MergeSort( $x_1, \dots, x_{\lfloor \frac{n}{2} \rfloor}$ ), MergeSort( $x_{\lfloor \frac{n}{2} \rfloor + 1}, \dots, x_n$ ))

```

Where the Merge function is a linear time algorithm combining two sorted lists (by comparing the first element in each list and moving the smaller element of

the two into our new list.)

MergeSort is correct!

- It calls itself on smaller instances until the division process terminates when it reaches a base case where each list has size 1
- MergeSort works trivially on the base cases
- This sort of strong induction-type proof works for just about all divide and conquer algorithms

MergeSort is efficient!

- The recurrence relation that we can construct to model the running time of MergeSort is given by

$$T(n) = 2 \cdot T\left(\frac{n}{2}\right) + O(n)$$

as we break the problem into two sub-problems of size $\frac{n}{2}$, plus a linear time sub-routine to combine the solutions (Merge.)

- This recurrence relation is in the proper form to use the Master Theorem and indeed we're in its Case II.
- By the master theorem the running time of MergeSort is $O(n \cdot \log n)$.
- The running time can also be proved by unwinding the recurrence.

2.2 Binary Search

Binary Search is an algorithm to search for whether a key k exists in a given sorted list.

```
BinarySearch( $a_1, a_2, a_n; k$ )
  while  $n > 0$ 
    if  $a_{\lfloor \frac{n}{2} \rfloor} = k$ 
      return true
    if  $a_{\lfloor \frac{n}{2} \rfloor} > k$ 
      return BinarySearch( $a_1, \dots, a_{\lceil \frac{n}{2} \rceil - 1}; k$ )
    if  $a_{\lfloor \frac{n}{2} \rfloor} < k$ 
      return BinarySearch( $a_{\lceil \frac{n}{2} \rceil + 1}, \dots, a_n; k$ )
  return false
```

Binary Search works!

- It works for the same strong induction argument as MergeSort

Binary Search is efficient!

- The recurrence relation that we can construct to model the running time of Binary Search is given by

$$T(n) = T\left(\frac{n}{2}\right) + O(1)$$

as we break the problem into one sub-problem of size $\frac{n}{2}$.

- This recurrence relation is case II of the master theorem.
- By the master theorem the running time of MergeSort is $O(\log n)$.
- Again, the running time can also be proved by unwinding the recurrence.

2.3 Fast Multiplication

2.4 Fast Matrix Multiplication

2.5 The Selection Problem

Suppose we're given a list of number and we want to find the k^{th} smallest number in that list. We can use the following divide and conquer algorithm:

```
Select( $\mathcal{S}, k$ )
  if  $|\mathcal{S}| = 1$ 
    return  $x_1$ 
  else
    set  $\mathcal{S}_L := \{x_i \in \mathcal{S} : x_i < x_1\}$ 
    set  $\mathcal{S}_R := \{x_i \in \mathcal{S} : x_i \geq x_1\}$ 

    if  $|\mathcal{S}_L| = k - 1$ 
      return  $x_1$ 
    else if  $|\mathcal{S}_L| > k - 1$ 
      return Select( $\mathcal{S}_L, k$ )
    else if  $|\mathcal{S}_L| < k - 1$ 
      return Select( $\mathcal{S}_R, k - 1 - |\mathcal{S}_L|$ )
```

Well.. this actually sucks because it runs in $\Omega(n^2)$ time. Note that it would actually be faster to just sort the list and then pull the k^{th} element (this can of course be done in $O(n \cdot \log n)$ time.) One idea that we can exploit to improve this is to use a random pivot rather than deterministically selecting x_1 as a pivot.

We will define a *good pivot* to be such neither \mathcal{S}_L , nor \mathcal{S}_R contain more than $\frac{3}{4}$ of \mathcal{S} , and a *bad pivot* to be otherwise. Note that by this definition the probability

that a pivot will be either good or bad (respectively) is 0.5.

For randomized algorithms we're always interested in the *expected runtime* of the algorithm rather than the worst case runtime. The recurrence relation, $\bar{T}(n)$, for the expected runtime for the randomized version of the selection is then given by

$$\bar{T}(n) \leq \frac{1}{2} \cdot \bar{T}\left(\frac{3n}{4}\right) + \frac{1}{2} \cdot \bar{T}(n) + O(n)$$

where the $\frac{1}{2} \cdot \bar{T}\left(\frac{3n}{4}\right)$ term represents the recursive call given a good pivot, and the $\frac{1}{2} \cdot \bar{T}(n)$ term represents the recursive call on a bad pivot. It's less than or equal to because we will only ever use one of those terms at once but not both. We can clean this recurrence up a bit to find the expected runtime.

$$\begin{aligned} \bar{T}(n) &\leq \frac{1}{2} \cdot \bar{T}\left(\frac{3n}{4}\right) + \frac{1}{2} \cdot \bar{T}(n) + O(n) \\ \frac{1}{2} \cdot \bar{T}(n) &\leq \frac{1}{2} \cdot \bar{T}\left(\frac{3n}{4}\right) + O(n) \\ \bar{T}(n) &\leq \bar{T}\left(\frac{3n}{4}\right) + O(n) \end{aligned}$$

Applying the master theorem to that recurrence gives $\bar{T}(n) = O(n)$

This is good but let's try to do better, i.e. let's see if we can find a deterministic linear time algorithm. It turns out we can but it requires a little trick called the *median of medians*.

It's best to illustrate this concept with an example. Suppose we have a list, $\mathcal{S} = \{x_1, x_2, \dots, x_n\}$, from which we want to find the median of medians. To do this we partition the list into $\lceil \frac{n}{5} \rceil$ groups of cardinality 5. Denote each such group by $G_1, \dots, G_{\lceil \frac{n}{5} \rceil}$. Then we'll sort each group and let z_i be the median of each group G_i . Let m be the median of the set $\mathcal{Z} = \{z_1, \dots, z_{\lceil \frac{n}{5} \rceil}\}$. Predictably, we call m the *median of medians*. This is significant because m will always be a good pivot. This is because there will always be at least $\frac{3n}{10} - 1$ numbers in \mathcal{S} less than m , and $\frac{3n}{10} - 1$ numbers in \mathcal{S} at least as large as m . Therefore $|\mathcal{S}_R| \leq \frac{7n}{10}$, and $|\mathcal{S}_L| \leq \frac{7n}{10}$.

Now that we have a deterministic way to guarantee we have a good pivot, all that's left is to bake it into a selection algorithm. Luckily, our algorithm won't vary too much since all we have to do is add a few steps at the beginning.

```

DetSelect( $\mathcal{S}, k$ )
  if  $|\mathcal{S}| = 1$ 
    return  $x_1$ 
  else
    partition  $\mathcal{S}$  into  $\lceil \frac{n}{5} \rceil$  groups of 5
    for  $j$  in range  $\left(1, \lceil \frac{n}{5} \rceil\right)$ 
      let  $z_j$  be the median of the group  $G_j$ .
    let  $\mathcal{Z} = \{z_1, \dots, z_{\lceil \frac{n}{5} \rceil}\}$ 
    set  $m := \text{DetSelect}\left(\mathcal{Z}, \lceil \frac{n}{10} \rceil\right)$ 

    set  $\mathcal{S}_L := \{x_i \in \mathcal{S} : x_i < m\}$ 
    set  $\mathcal{S}_R := \{x_i \in \mathcal{S} : x_i \geq m\}$ 

    if  $|\mathcal{S}_L| = k - 1$ 
      return  $m$ 
    if  $|\mathcal{S}_L| > k - 1$ 
      return DetSelect( $\mathcal{S}_L, k$ )
    if  $|\mathcal{S}_L| < k - 1$ 
      return DetSelect( $\mathcal{S}_R, k - 1 - |\mathcal{S}_L|$ )

```

The recursive formula for the running time of this algorithm is

$$T(n) \leq T\left(\frac{7n}{10}\right) + T\left(\frac{n}{5}\right) + O(n)$$

where the $T\left(\frac{n}{5}\right)$ term is the time to find the median of medians, the $T\left(\frac{7n}{10}\right)$ term is pivoting based on the median of medians. Unfortunately, this doesn't work with the master theorem and we can't simplify it down to a form that does (unlike with the randomized algorithm, we need both of the $T(\dots)$ terms.) Since it doesn't work with the master theorem, we need to use the recursion tree method. From the recursion tree method, $T(n) = O(n)$. So there it is.. a linear time deterministic algorithm to select the k^{th} smallest element from a list!

3 Graph Algorithms

4 Greedy Algorithms

5 Dynamic Programming

6 Network Flows

7 Data Structures